

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
insurance_dataset = pd.read_csv('/content/insurance.csv')
```

```
insurance_dataset.head()
```

```
↗
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
insurance_dataset.shape
```

```
↗ (1338, 7)
```

```
insurance_dataset.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Categorical Features:

- Sex
- Smoker
- Region

```
# checking for missing values
insurance_dataset.isnull().sum()
```

```
↗
```

	0
age	0
sex	0
bmi	0
children	0
smoker	0
region	0
charges	0

Data Analysis

```
# statistical Measures of the dataset
insurance_dataset.describe()
```

	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010

### Encoding the categorical features

```
# encoding sex column
insurance_dataset.replace({'sex':{'male':0,'female':1}}, inplace=True)
```

```
3 # encoding 'smoker' column
insurance_dataset.replace({'smoker':{'yes':0,'no':1}}, inplace=True)
```

```
# encoding 'region' column
insurance_dataset.replace({'region':{'southeast':0,'southwest':1,'northeast':2,'northwest':3}}, inplace=True)
```

```
<ipython-input-8-7d5826986d65>:2: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future ver
insurance_dataset.replace({'sex':{'male':0,'female':1}}, inplace=True)
<ipython-input-8-7d5826986d65>:5: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future ver
insurance_dataset.replace({'smoker':{'yes':0,'no':1}}, inplace=True)
<ipython-input-8-7d5826986d65>:8: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future ver
insurance_dataset.replace({'region':{'southeast':0,'southwest':1,'northeast':2,'northwest':3}}, inplace=True)
```

```
insurance_dataset.head()
```

	age	sex	bmi	children	smoker	region	charges
<b>0</b>	19	1	27.900	0	0	1	16884.92400
<b>1</b>	18	0	33.770	1	1	0	1725.55230
<b>2</b>	28	0	33.000	3	1	0	4449.46200
<b>3</b>	33	0	22.705	0	1	3	21984.47061
<b>4</b>	32	0	28.880	0	1	3	3866.85520

### Splitting the Features and Target

```
X = insurance_dataset.drop(columns='charges', axis=1)
Y = insurance_dataset['charges']
```

```
print(X)
```

```

age  sex    bmi  children  smoker  region
0    19    1  27.900         0       0       1
1    18    0  33.770         1       1       0
2    28    0  33.000         3       1       0
3    33    0  22.705         0       1       3
4    32    0  28.880         0       1       3
...    ...    ...    ...    ...    ...    ...
1333  50    0  30.970         3       1       3
1334  18    1  31.920         0       1       2
1335  18    1  36.850         0       1       0
1336  21    1  25.800         0       1       1
1337  61    1  29.070         0       0       3
```

```
[1338 rows x 6 columns]
```

```
print(Y)
```

```

0    16884.92400
1     1725.55230
2     4449.46200
3    21984.47061
4     3866.85520
...
1333  10600.54830
```

```

1334      2205.98080
1335      1629.83350
1336      2007.94500
1337      29141.36030
Name: charges, Length: 1338, dtype: float64

```

### Splitting the data into Training data & Testing Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(1338, 6) (1070, 6) (268, 6)
```

### Model Training

```
# loading the Linear Regression model
regressor = LinearRegression()
```

```
regressor.fit(X_train, Y_train)
```

```

LinearRegression
LinearRegression()

```

### Model Evaluation

```
# prediction on training data
training_data_prediction = regressor.predict(X_train)
```

```
# R squared value
r2_train = metrics.r2_score(Y_train, training_data_prediction)
print('R squared vale : ', r2_train)
```

```
R squared vale : 0.751505643411174
```

```
# prediction on test data
test_data_prediction = regressor.predict(X_test)
```

```
# R squared value
r2_test = metrics.r2_score(Y_test, test_data_prediction)
print('R squared vale : ', r2_test)
```

```
R squared vale : 0.7447273869684076
```

```
input_data = (31,1,25.74,1,1,0)
```

```
# changing input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
# reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
print(input_data_reshaped)
prediction = regressor.predict(input_data_reshaped)
print(prediction)
```

```
print('The insurance cost is USD ', prediction[0])
```

```

[[31.  1.  25.74  1.  1.  0. ]]
[4340.35495946]
The insurance cost is USD 4340.354959456534
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but LinearRegression
warnings.warn(

```

### Training the Random Forest Model

```
# Assuming 'charges' is the target variable and others are features
X = insurance_dataset.drop('charges', axis=1) # Features (all columns except target)
y = insurance_dataset['charges'] # Target (the 'charges' column)
```

```
# Handling categorical variables using one-hot encoding
X = pd.get_dummies(X, drop_first=True) # Drop the first category to avoid multicollinearity
```

```
# Checking the preprocessed features
X.head()
```

```
↗
```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3

### Splitting the data

```
# Splitting the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Checking the shape of the splits to verify
print(f"Training data shape: X_train = {X_train.shape}, y_train = {y_train.shape}")
print(f"Testing data shape: X_test = {X_test.shape}, y_test = {y_test.shape}")
```

```
↗ Training data shape: X_train = (1070, 6), y_train = (1070,)
Testing data shape: X_test = (268, 6), y_test = (268,)
```

```
from sklearn.ensemble import RandomForestRegressor # For regression tasks
```

```
# Creating a RandomForestRegressor model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
```

```
# Training the model on the training data
rf_model.fit(X_train, y_train)
```

```
↗
```

RandomForestRegressor ⓘ ?  
RandomForestRegressor(random\_state=42)

### Making Predictions

```
# Predicting the target variable on the test data
y_pred = rf_model.predict(X_test)
```

```
# Showing the first few predictions to get a sense of the output
y_pred[:10]
```

```
↗ array([ 9964.4411712,  5614.908105 , 28122.751816 , 12317.3170911,
        34592.244544 ,  8330.8161489,  2185.8946415, 14516.191212 ,
        5719.9984878, 10166.264253 ])
```

### Evaluate the model

```
# Calculating the Mean Squared Error and Root Mean Squared Error
mse = metrics.mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
```

```
# Printing the evaluation metrics
print(f"Mean Squared Error: {mse}")
print(f"Root Mean Squared Error: {rmse}")
```

```
↗ Mean Squared Error: 20605006.150018733
Root Mean Squared Error: 4539.273746979657
```

```
# Getting the feature importances from the trained model
feature_importances = rf_model.feature_importances_
```

```
# Creating a DataFrame to display feature names and their importance
importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importances
})
```

```
# Sorting the features by importance in descending order
importance_df = importance_df.sort_values(by='Importance', ascending=False)
```

```
# Displaying the feature importance
print(importance_df)
```

```
↗
```

	Feature	Importance
4	smoker	0.608618
2	bmi	0.216403
0	age	0.134356
3	children	0.019627
5	region	0.014326
1	sex	0.006670

```
# Assuming the trained RandomForest model is stored in 'rf_model'
input_data = (31, 1, 25.74, 1, 1, 0) # Example input data
```

```
# Converting the input_data into a numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
# Reshaping the array to match the model input format (1 row, multiple columns)
input_data_resaped = input_data_as_numpy_array.reshape(1, -1)
```

```
# Printing the reshaped input data for reference
print(input_data_resaped)
```

```
# Using the trained Random Forest model to make a prediction
prediction = rf_model.predict(input_data_resaped)
```

```
# Printing the predicted insurance cost
print(f'The predicted insurance cost is USD {prediction[0]}')
```

```
↗ [[31.  1. 25.74  1.  1.  0.  ]]
The predicted insurance cost is USD 4878.8846039
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but RandomForestRegressor has feature names
  warnings.warn(
```

```
import pickle
import joblib
filename='InsuranceCostPredictor.pkl'
joblib.dump(rf_model, filename)
```

```
↗ ['InsuranceCostPredictor.pkl']
```